

Deep Learning for NLP 2020

Exercise 09

June 14, 2020

1 Pingo

Try to find the right answer(s) to each question on your own or in a group with your colleagues. The interactive survey will be conducted near the end of the practice class.

- Which of the following statements about CNNs are correct?
 - ☐ The convolution operation creates an activation map (aka feature map) for every filter.
 - ☐ In CNNs, one usually applies a pooling layer first, followed by a convolutional layer.
 - ☐ A stride of 2 means that every convolution operation folds two neighboring input values.
 - ☐ When applying CNNs on NLP tasks, the features learned by every filter matrix can be easily visualized and interpreted.
- Which facts about regularization strategies are true?
 - ☐ Dropout can be applied for CNNs.
 - ☐ The dropout probability p is a hyperparameter.
 - ☐ Early stopping requires a validation set.

2 Theoretic Background of CNNs

1. Is a convolution over individual words (window size $k = 0$, i.e. one word per convolution window) useful? Explain in up to three sentences.

3 Dimensions and Parameters

We use a convolutional neural network for sentence classification.

Each input sentence consists of 197 words. Each word is represented by a vector from a 300-dimensional embedding space. There are 5581 unique words. The network consists of a trainable embedding layer followed by a convolutional layer, a global max-pooling layer and another fully-connected hidden layer. 111 filters, each with a window size of $k = 2$ convolve over the input. The stride is 1. The convolution is narrow. The hidden layer has 42 neurons and uses dropout with keep probability $p = 80\%$. The output layer is a single neuron with sigmoid activation function.

Using pen and paper...

-
1. Compute the output shape of each layer.
 2. Compute the number of trainable parameters of each layer. Don't forget the bias for the filters and the hidden layers!